

OmniMed-FL: A Robust Multimodal Federated Learning Framework for Clinical Diagnosis

¹Ayush Debnath, ^{2,3,4}Ruelia Saha, ²Sudip Misra

¹Department of Mechanical Engineering

Indian Institute of Technology Kharagpur, India-721302

²Computer Science & Engineering Department

IIT (Indian Institute of Technology) Kharagpur, India-721302

³KTH Royal Institute of Technology, Stockholm, Sweden

⁴SRM University- AP, India

Email: {ayush.d@kgpian.iitkgp.ac.in, ruelia.saha.rs@gmail.com, sudipm@iitkgp.ac.in}

Abstract—We present OmniMed-FL, a controlled multimodal federated-learning systems study for five-class image–text classification. The proxy corpus contains 2,400 public chest radiographs, 600 synthetic COVID-19 image surrogates, and 3,000 class-conditioned synthetic notes paired to images by class rather than patient. Public DistilBERT and ViT-Base/16 encoders are fine-tuned with random task modules; no pooled-client model initializes the operational runs. We vary Dirichlet skew, client count, anti-collapse components, initialization, eight operational fusion rules, missing text, and matched federated controls, while reporting runtime, peak allocated GPU memory, and formula-derived FP32 traffic. At $K=5$ and severe skew ($\alpha=0.1$), local-only training obtains 0.310 Macro F1 (five-client mean), versus 0.652 ± 0.057 for FedAvg, 0.685 ± 0.064 for FedProx, and 0.192 ± 0.129 for our SCAFFOLD-style AdamW adaptation; the 0.033 FedProx–FedAvg gap is smaller than either two-seed sample standard deviation. In the heterogeneity sweep, FedAvg changes from 0.896 ± 0.004 at $\alpha=5$ to 0.674 ± 0.185 at $\alpha=0.1$. A branch audit records 0.248/0.324/0.573 GiB FP32 payload per client per round for text/image/multimodal models. Because the notes and one image class are synthetic, pairing is not patient-level, and no external clinical validation is performed, these are descriptive proxy comparisons rather than estimates of diagnostic performance, privacy, or deployment readiness.

Index Terms—Multimodal Federated Learning, Non-IID Data, Multimodal Fusion, Clinical Classification, Communication Cost, Retrieval

I. INTRODUCTION

CLINICAL prediction often draws on both images and textual context. Centralized image models [1], [2] and medical vision–language models can combine these signals, but cross-institutional training would ordinarily require data to leave its source. Federated learning (FL) instead exchanges parameter updates while retaining each client’s raw examples [3]. This data-locality property is useful, although FL alone is not a privacy guarantee and does not remove risks from model updates.

Federated medical imaging [4]–[6] and centralized medical vision–language modeling have established precedents, with recent work also studying their intersection. The combination still raises questions that a single favorable run cannot answer: how performance changes with client skew, whether anti-collapse components help, whether fusion rankings survive seed

variation, and what additional memory and communication the multimodal branch incurs.

We study these questions through OmniMed-FL, a controlled five-class proxy benchmark. Its public radiographs and synthetic clinical-style notes are not patient-matched data, so the work evaluates the learning system rather than clinical diagnosis. This scope permits internally matched comparisons while making the boundary between a systems result and a clinical claim explicit. Here, “robust” denotes stress testing across partitions, client counts, and missing modalities, not certified clinical robustness.

A. Contribution

The main *contributions* of this paper are as follows:

- We implement a controlled multimodal FL protocol that begins from public pretrained encoders and random task modules, never from a model trained on pooled client examples.
- We evaluate heterogeneity, client count, anti-collapse components, initialization, and eight fusion rules under one multi-round FL protocol. Separate matched adaptations cover FedMME’s one-shot and P-FIN’s missing-text settings.
- We report realized client skew, wall-clock time, peak allocated GPU memory and FP32 communication payload, and evaluate retrieval over the full held-out split rather than a five-query demonstration.

II. RELATED WORK

Centralized radiograph classifiers [1], [2] and Vision Transformers [7], [8] provide established image-only precedents. The clinical FL studies cited here are also image-only: Sheller *et al.* evaluate federated averaging, Dou *et al.* conduct multinational external validation, and Kaissis *et al.* combine differential privacy, secure aggregation, and encrypted inference [4]–[6]. FedProx and SCAFFOLD address heterogeneous-client optimization through proximal regularization and control variates, respectively [9], [10].

BiomedCLIP and MedCLIP learn medical image–text representations, while LLaVA-Med instruction-tunes a biomedical




Fig. 1. OmniMed-FL architecture. DistilBERT and ViT-Base/16 encode the two proxy modalities; lightweight multimodal fusion combines them; FedAvg coordinates simulated clients without exchanging their raw examples.

vision–language assistant; all three are centralized [11]–[13]. Recent multimodal FL includes FedMME’s one-shot hard-voting ensemble and P-FIN’s uncertainty-aware feature imputation for missing modalities [14], [15]. Because their published datasets, backbones, tasks, and budgets differ, we report clearly labeled adaptations on our common corpus and local-epoch budget rather than compare cross-paper scores or claim an exact reproduction.

III. SYSTEM DESIGN

A. Architecture

Figure 1 shows the OmniMed-FL pipeline. Clinical-style notes are tokenized with WordPiece to 128 tokens; chest images are resized to 224×224 and ImageNet-normalized. DistilBERT’s [16] [CLS] representation is projected to $\mathbf{h}_t \in \mathbb{R}^{256}$, while the mean-pooled ViT token sequence is projected to $\mathbf{h}_v \in \mathbb{R}^{512}$.

The two representations are combined by eight lightweight operators: projected concatenation, residual cross-attention, gated residual fusion, dual-projection concatenation, one-layer decoder A, a distinct two-query/two-layer decoder B, non-residual 384-dimensional cross-attention, and a two-token Transformer encoder. Historical code keys resemble foundation-model names, but these are compact fusion rules rather than implementations of those models. Each feeds the same two-layer five-class MLP, and all eight are compared inside the operational FL loop with the same data, partition, optimizer, and local budget.

B. Federated and Local Objectives

Let K clients hold disjoint shards $\mathcal{D}_k$ of the training split, with $n_k = |\mathcal{D}_k|$ and $n = \sum_k n_k$, and let θ collect the trainable parameters on the fusion path: both encoders, the projection

layers, the fusion operator and the classification head. The federated objective is the sample-weighted empirical risk

$$\min_{\theta} F(\theta) = \sum_{k=1}^K \frac{n_k}{n} F_k(\theta), \quad (1)$$

$$F_k(\theta) = \frac{1}{n_k} \sum_{(x,z,y) \in \mathcal{D}_k} \ell(f_{\theta}(x,z), y).$$

where x is a radiograph, z its class-paired note and f_{θ} the fused five-class classifier. Shards are never exchanged; the server receives only the trainable tensors enumerated above.

The local loss combines a focal term with a batch-level diversity penalty,

$$\ell = (1 - p_y)^{\gamma} \text{CE}_{\epsilon}(f_{\theta}(x,z), y) + \lambda \left(1 - \frac{H(\bar{p})}{\log C}\right), \quad (2)$$

with focusing parameter $\gamma=2$, label smoothing $\epsilon=0.05$, $C=5$ classes, and $\bar{p}$ the mean predicted distribution over the mini-batch. The second term grows when a client’s predictions concentrate on few classes, a behavior that can occur under severe Dirichlet skew. Setting $\lambda=0$ removes it; together with a class-balanced batch sampler it forms the anti-collapse stack ablated in Section IV.

At each round, the server broadcasts $\theta^{(t)}$; each eligible client performs three class-balanced local AdamW epochs using (2), clips the gradient norm at one, and returns its trainable tensors. The server applies the sample-weighted update in (1). Operational runs begin with public DistilBERT and ViT-Base/16 encoder weights plus random projection, fusion, and classification heads; no pooled-client model initializes them. Both encoders are fine-tuned, whereas unused unimodal heads and the inactive CNN fallback are excluded from optimization and communication.

C. Controlled Proxy Corpus

Image data. We form a balanced 3,000-image proxy corpus (600 examples per class). The public radiographs come from the Kermany pneumonia collection [17] and the NIH ChestX-ray dataset [18]. Loader logs verify 2,400 public images covering Normal, Pneumonia, Pleural Effusion and Cardiomegaly. The originally specified COVID-19 source was unavailable at execution time, so the 600 COVID-19 images are synthetic X-ray-style surrogates with bilateral opacity patterns. Images are resized to 224×224 , normalized with ImageNet statistics, shuffled deterministically and split 80/20 into 2,400 training and 600 validation examples.

Clinical text data. No patient EHR data are used. Each image label is paired at the class level with a synthetic clinical-style note generated from observation, symptom, context and indicator slots. Twenty-five percent of notes draw all slots from the target class, 35% mix target and non-target slots, and 40% are heavily mixed while retaining a probabilistic target-class indicator. This controlled construction reduces direct label wording, but the image and note are not from the same patient and must not be described as a real image–report pair.

What this corpus does and does not support. Every experimental arm sees the same controlled corpus, which supports

internally controlled, descriptive comparisons between federated algorithms, fusion rules and ablations. The absolute scores are not diagnostic-performance estimates: class-conditioned templates can create shortcuts, label-level pairing removes real patient discordance and missingness, and the synthetic COVID-19 surrogates may not resemble clinical COVID-19 radiographs. The split is also not patient-grouped across source institutions. We therefore make no claim about clinical utility, deployment readiness, or generalization to real paired reports.

D. Experimental Setup

Hyperparameters. Each client uses AdamW with learning rate 10^{-4} and weight decay 0.01 for three local epochs per round. The reviewer suite uses eight rounds, batch size 16, FP32 tensors without automatic mixed precision, and diversity-loss weight 1.0. The reference setting is five clients with Dirichlet $\alpha = 1.0$; Section IV varies α and the number of clients explicitly. Seeds 0 and 1 jointly control task-head initialization and the realized client partition. Experiments run in PyTorch on one NVIDIA H100 NVL device in a server shared with unrelated workloads. The GPU-synchronized timer covers sequential local training and aggregation within each round; validation, setup, I/O, and networking are excluded, while server contention remains uncontrolled. CUDA deterministic algorithms are not forced, so seeds control software RNG streams without guaranteeing bitwise-identical reruns. Simulated clients train sequentially on one GPU, so time is an execution-cost record rather than distributed throughput. Table I records the fixed protocol.

Partitioning. For each class, a Dirichlet draw allocates training indices across clients. Smaller α creates more concentrated client label distributions; $\alpha \rightarrow \infty$ approaches IID allocation [19]. Because nominal α does not fully describe a finite realized split, we also report the mean client share occupied by its dominant class. The same split is held fixed across rounds for a given seed. Shards with fewer than four training examples are not locally updated; scalability results therefore report realized active and nominal K .

Matched recent-method controls. Our matched FedMME-style arm uses 24 local epochs rather than FedMME’s native 100, uploads each client model once, and applies its equal-weight hard voting; our implementation breaks ties by mean softmax [14]. This preserves the common corpus, shards, encoders, optimizer, and matched local compute, but substitutes our five-class inputs for FedMME’s generated-report pipeline. The P-FIN-style stress test fixes three multimodal and two image-only clients, projects both encoders to 256 normalized features, and compares zero fill, deterministic FIN, Gaussian beta-NLL imputation ($\beta=0.5$), and uncertainty-weighted aggregation with Fed-UQ-Avg balance weight 0.6 and temperature 0.2 [15]. It uses the common concat classifier rather than P-FIN’s native fusion stack. FedProx uses $\mu=0.01$. These are matched adaptations, not exact reproductions; their native scopes and communication schedules remain different.

Rationale for Table I. Within each ablation or comparison family, only the named factor changes. Three local epochs and

TABLE I
COMMON SIMULATION PARAMETERS.

Parameter	Value
Local Epochs (E)	3
Federated Rounds (T)	8
Batch Size	16
AdamW Learning Rate	1×10^{-4}
Weight Decay	0.01
Diversity Reg. (λ)	1.0
Systems-grid K (reference)	3, 5, 10, 20 (5)
Systems-grid α (reference)	0.1, 1, 5 (1)
$K=5$ additional α	0.3, 0.5
Seeds	0, 1
Initialization	Public encoders + random task modules
Precision	FP32 tensors (no AMP)

eight rounds give a tractable 24-local-epoch budget, while batch size 16 holds both encoders and one local model copy in GPU memory. The Cartesian systems grid crosses $K=\{3, 5, 10, 20\}$ with $\alpha=\{0.1, 1, 5\}$; the $K=5$ heterogeneity sweep additionally includes $\alpha=0.3$ and 0.5. Thus the study spans severe to near-IID label skew, with $K=5$, $\alpha=1$ as its reference. Two seeds were computationally feasible and support descriptive, not significance, claims. The 10^{-4} rate and $\lambda=1$ are fixed implemented defaults, not per-model optima.

IV. RESULTS

A. Matched Performance and Ablations

At $\alpha=0.1$, local-only training averages 0.310 Macro F1 across the five clients in one realized partition. Under matched two-seed partitions, OmniMed-FL with FedAvg reaches 0.652 ± 0.057 and FedProx reaches 0.685 ± 0.064 (Fig. 2(a)). Their 0.033 difference is smaller than either sample standard deviation. SCAFFOLD-AdamW obtains 0.192 ± 0.129 and oscillates in Fig. 2(b); this identifies instability in our adaptive-optimizer control-variate adaptation and should not be generalized to classical SGD SCAFFOLD. In the FedMME-style matched adaptation, five independent client models train for 24 epochs and their held-out predictions are combined by equal hard plurality with mean-softmax tie breaking. It preserves the one-shot ensemble idea but is not an exact FedMME reproduction.

The anti-collapse effects are partition-dependent (Fig. 2(c)). At $\alpha=0.1$, the complete balanced-loader/diversity-loss stack gives 0.716 ± 0.014 F1. Removing the balanced loader gives 0.585 ± 0.148 , removing the diversity term gives 0.701 ± 0.011 , and removing both gives 0.657 ± 0.057 ; the overlaid diversity markers distinguish transient class collapse from lower F1. At $\alpha=1$, no component has a resolved two-seed benefit. Independent baseline, heterogeneity-sweep, and ablation executions give 0.652 ± 0.057 , 0.674 ± 0.185 , and 0.716 ± 0.014 for the nominally repeated severe-skew FedAvg/concat setting; all are retained rather than pooled or selected. The eight fusion rules in Fig. 2(d) are trained inside the same operational FL loop rather than screened on pooled data.

Initialization is separated from early stopping in Fig. 2(e). All-random uses the same DistilBERT/ViT architectures with random encoder and task weights. After initialization, all three arms use the same eight-round, three-local-epoch FL schedule.




Fig. 2. Matched performance and ablations. (a) Severe-skew baselines use the same corpus and partition family; local-only is the five-client mean from one partition, whereas federated whiskers are sample standard deviations over two seeds. FedMME-style is a one-shot hard-voting ensemble after each client trains independently for 24 local epochs; iterative arms use eight rounds of three epochs. (b) Iterative convergence. (c) Anti-collapse F1; diamonds are the minimum predicted-class diversity during training. (d) Eight operational fusion rules. (e) Federated initialization, with round-1 diamonds, beside a visually separated 12-epoch pooled-data abort control that is non-federated. (f) P-FIN-style adaptation with three multimodal and two image-only clients; the primary metric forces text missing for every held-out example, while full-text F1 is secondary. Except where noted, bars are two-seed means with sample-SD whiskers.

The pooled-start arm additionally reuses one seed-0 checkpoint selected from a 12-epoch centralized pooled-data run for both federated seeds; this extra pretraining is excluded from the 8×3 budget. It begins at higher round-1 F1 (0.843 versus 0.704 for public encoders) but finishes lower (0.818 versus 0.839). Only this labeled ablation uses the pooled checkpoint; operational fusion, grid, and baseline arms start from public encoders with random task modules. The neighboring abort-off/on bars (0.845 ± 0.013 and 0.884 ± 0.003) are a separate centralized 12-epoch diagnostic. Both settings complete all 12 epochs with minimum predicted-class diversity 1.0, so the collapse-triggered learning-rate boost/abort branch never activates; their difference is independent non-bitwise-deterministic rerun variation, not an estimated early-abort effect. Figure 2(f) tests missing text with 256-dimensional L_2 -normalized features and concat+MLP fusion in place of P-FIN’s bidirectional attention. It is a matched adaptation, not an exact reproduction; forced missing text is primary and full-text performance is secondary.

B. Scalability, Communication, and Resource Cost

Figure 3(a) crosses $K=3, 5, 10, 20$ with $\alpha=0.1, 1, 5$. At $K=5$, F1 changes from 0.896 ± 0.004 near IID ($\alpha=5$) to 0.674 ± 0.185 under severe skew ($\alpha=0.1$); the additional $\alpha=0.3$ and 0.5 checks give 0.769 ± 0.043 and 0.824 ± 0.113 . The much wider severe-skew spread, together with the realized allocation in Fig. 3(f), shows why a nominal Dirichlet parameter and a single seed are insufficient. Cells report A/K (or the two-seed A range); the audited runner skips local optimization for realized shards smaller than four samples.

The timing and memory heat maps in Fig. 3(b)–(c) provide the requested systems measurements across client counts and data distributions. They time the GPU-synchronized sequential local-training/aggregation block and record peak allocated memory within that block on one H100 NVL shared with unrelated workloads; validation, setup, I/O, and networking are excluded. They are not multi-node-throughput measurements. For trainable-state payload $|\theta|b$ bytes, the idealized bidirectional traffic is

$$V_{\text{nom}} = 2KT|\theta|b, \quad (3)$$

Here $|\theta|=153,935,621$, $b=4$ bytes, and $T=8$, so the 615,742,484-byte payload gives 27.5/45.9/91.8/183.5 GiB for $K=3/5/10/20$, using 1 GiB = 2^{30} bytes. This nominal accounting remains $2KT|\theta|b$ when $A < K$; substituting A would be an alternative active-client estimate, not a measurement made here. It excludes serialization, headers, retransmission, secure aggregation, compression, and algorithm-specific state.

The single-run branch audit (Fig. 3(e)) makes the multimodal cost explicit. Text/image/multimodal F1 is 0.893/0.740/0.782; their FP32 payloads are 0.248/0.324/0.573 GiB per client per round, summed timed-round sections are 18.5/33.6/45.0 min, and peak allocated memory is 1.69/3.18/4.77 GiB. In this single proxy run, the text branch exceeds concatenation, so these data do not support a blanket claim that multimodal fusion is always superior; class-conditioned template shortcuts are one plausible explanation.

C. Retrieval Evidence

Exact inner-product FAISS search [20] over TF-IDF note vectors gives condition-macro averages of 0.522 top-1 same-label accuracy, 0.462 same-label precision@5, and 0.707 mean top-1 cosine similarity across the five condition rows of the 600-query evaluation (Fig. 4). This replaces the earlier five-query illustration. It is a synthetic-note retrieval check, not a fine-tuned-DistilBERT result, generated-answer evaluation, or evidence of clinical retrieval quality.

V. DISCUSSION AND LIMITATIONS

The experiments support systems-level conclusions only. Strong partition sensitivity persists despite anti-collapse measures: at $\alpha=0.1$, even identically configured FedAvg reruns vary substantially, and the audited runner skips local optimization for shards with fewer than four samples. The client-count study is a sequential simulation; it does not model concurrent workers, network latency, stragglers, intermittent dropout, or heterogeneous accelerators. Robust aggregation against poisoned or Byzantine updates is also untested. Accordingly, “robust” in the title denotes the reported stress tests, not a formal guarantee.

Communication is a second boundary. Equation (3) exposes model-state scaling, while Fig. 3(e) shows the larger multimodal




Fig. 3. Scalability and resource audit. Heat maps show two-seed mean $\pm$ sample SD for (a) Macro F1; A/K is the active/nominal count or a two-seed active-count range, (b) the GPU-synchronized round section, and (c) peak allocated GPU memory within that section. (d) Formula-derived two-way FP32 model-state volume counts one download and upload for every nominal client in each of eight rounds, including shards skipped from local optimization; it is not network traffic and has no stochastic error bar. (e) One-run branch costs at $\alpha=1$, $K=5$; time is the sum of timed-round sections and payload is per client per round. (f) The realized severe-skew seed-0 allocation; text is sample count and color is within-client class share.



Fig. 4. Exact TF-IDF/FAISS retrieval over all 600 held-out queries. Bars show top-1 same-condition-label accuracy and same-condition-label precision@5; diamonds show mean top-1 cosine similarity with standard-deviation whiskers. This evaluates retrieval, not generation.

payload and memory footprint. A deployment study must additionally measure end-to-end traffic and latency under compression, partial participation, failures, and secure aggregation. This simulation implements neither secure aggregation nor the differentially private training and privacy accounting of Abadi *et al.* [21]. It also does not evaluate attacks on model updates; retaining raw examples locally is therefore not a formal privacy guarantee.

Clinical interpretation is more restricted. All 3,000 notes and the 600-image COVID-19 class are synthetic; the other 2,400 images are public radiographs. Image-text pairing is by class rather than patient; class-conditioned templates can encode label shortcuts; and label-level pairing removes the discordance, missingness, temporal context, and site effects of real records. The split is not patient- or institution-grouped, and there is no external or prospective validation. Thus neither absolute F1 nor retrieval scores estimate diagnostic safety, generalization, clinical utility, or deployment readiness.

Finally, two seeded runs without enforced deterministic CUDA algorithms permit descriptive means and sample standard deviations, not significance claims. The matched FedMME- and P-FIN-style arms retain important method ideas under a common corpus and 24-local-epoch budget, but use different architectures/schedules and substitute our five-class inputs and fusion stack; they are not exact reproductions of the source papers. Larger multi-site datasets, more seeds, native-method replications, calibrated uncertainty, privacy/security tests, and

prospective evaluation are necessary next steps.

VI. CONCLUSION

OmniMed-FL is a controlled multimodal-FL systems study rather than a clinical validation. Reported performance and resource cost are conditional on partition realization and protocol choices; the fusion, initialization, and missing-modality comparisons are descriptive stress tests. Real patient-paired multi-institutional studies and deployment safeguards remain necessary before clinical use.

ACKNOWLEDGMENT

The work reported in this paper was done as part of the INAE Chair Professorship grant received by the third author.

REFERENCES

- [1] P. Rajpurkar, J. Irvin, K. Zhu, et al., “CheXNet: Radiologist-level pneumonia detection on chest X-rays with deep learning,” *arXiv:1711.05225*, 2017.
- [2] J. Irvin, P. Rajpurkar, M. Ko, et al., “CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison,” in *Proc. AAAI Conf. Artificial Intelligence*, vol. 33, no. 1, pp. 590–597, 2019.
- [3] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. 20th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, PMLR, vol. 54, pp. 1273–1282, 2017.
- [4] M. J. Sheller, B. Edwards, G. A. Reina, et al., “Federated learning in medicine: facilitating multi-institutional collaborations without sharing patient data,” *Scientific Reports*, vol. 10, art. 12598, 2020.
- [5] Q. Dou, T. Y. So, M. Jiang, et al., “Federated deep learning for detecting COVID-19 lung abnormalities in CT: a privacy-preserving multinational validation study,” *npj Digital Medicine*, vol. 4, art. 60, 2021.
- [6] G. Kaissis, A. Ziller, J. Passerat-Palmbach, et al., “End-to-end privacy preserving deep learning on multi-institutional medical imaging,” *Nature Machine Intelligence*, vol. 3, pp. 473–484, 2021.
- [7] A. Dosovitskiy, L. Beyer, A. Kolesnikov, et al., “An image is worth 16 $\times$ 16 words: Transformers for image recognition at scale,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2021.
- [8] Z. Liu, Y. Lin, Y. Cao, et al., “Swin Transformer: Hierarchical vision transformer using shifted windows,” in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2021, pp. 10012–10022.
- [9] T. Li, A. K. Sahu, M. Zaheer, et al., “Federated optimization in heterogeneous networks,” in *Proc. Mach. Learn. Syst. (MLSys)*, vol. 2, pp. 429–450, 2020.
- [10] S. P. Karimireddy, S. Kale, M. Mohri, et al., “SCAFFOLD: Stochastic controlled averaging for federated learning,” in *Proc. 37th Int. Conf. Machine Learning (ICML)*, PMLR, vol. 119, pp. 5132–5143, 2020.
- [11] S. Zhang, Y. Xu, N. Usuyama, et al., “BiomedCLIP: A multimodal biomedical foundation model pretrained from fifteen million scientific image-text pairs,” *arXiv:2303.00915*, 2023.
- [12] C. Li, C. Wong, S. Zhang, et al., “LLaVA-Med: Training a large language-and-vision assistant for biomedicine in one day,” in *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, vol. 36, pp. 28541–28564, 2023.
- [13] Z. Wang, Z. Wu, D. Agarwal, and J. Sun, “MedCLIP: Contrastive learning from unpaired medical images and text,” in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2022, pp. 3876–3887.

- [14] N. Wang, Y. Deng, S. Fan, J. Yin, and S.-K. Ng, "Multi-modal one-shot federated ensemble learning for medical data with vision large language model," *arXiv:2501.03292*, 2025.
- [15] N. F. Shahid, M. Ahmed, M. A. Haider, et al., "Probabilistic feature imputation and uncertainty-aware multimodal federated aggregation," in *Proc. Medical Imaging with Deep Learning (MIDL)*, PMLR, vol. 315, pp. 3593–3607, 2026.
- [16] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, "DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter," *arXiv:1910.01108*, 2019.
- [17] D. S. Kermany, M. Goldbaum, W. Cai, et al., "Identifying medical diagnoses and treatable diseases by image-based deep learning," *Cell*, vol. 172, no. 5, pp. 1122–1131.e9, 2018.
- [18] X. Wang, Y. Peng, L. Lu, et al., "ChestX-ray8: Hospital-scale chest X-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 2097–2106.
- [19] T.-M. H. Hsu, H. Qi, and M. Brown, "Measuring the effects of non-identical data distribution for federated visual classification," *arXiv:1909.06335*, 2019.
- [20] J. Johnson, M. Douze, and H. Jégou, "Billion-scale similarity search with GPUs," *IEEE Trans. Big Data*, vol. 7, no. 3, pp. 535–547, 2021.
- [21] M. Abadi, A. Chu, I. Goodfellow, et al., "Deep learning with differential privacy," in *Proc. ACM SIGSAC Conf. Computer and Communications Security (CCS)*, 2016, pp. 308–318.